

Question 1: Pipelining

Stages: 5 stages, taking times 3ns, 6ns, 1ns, 10ns, and 5ns.

a. No Pipelining:

If the stages do not operate in parallel, then one instruction at a time falls through each stage, finishing with the last step before the next instruction is executed. Each and every instruction, then, takes $(3 + 6 + 1 + 10 + 5)$ ns of time, so the total number of instructions executed per second is $1s / (3 + 6 + 1 + 10 + 5)$ ns, or $1\,000\,000\,000 / 25$ ns = 40 000 000 instructions per second, total.

b. With Pipelining.

If the stages operate in parallel, the issue is that each stage still depends on the results of the prior stage to operate on. From the perspective of any one instruction, it must still fall through each stage, and take 25ns as a result. However, after an instruction completes a stage, the next instruction is loaded into that stage, as opposed to waiting for the instruction to finish, even as the current instruction continues on to the next stage. From the perspective of processed stages (the end of the pipeline), once the first instruction is fully processed, the next instructions will be processed as quickly as the slowest part of the pipeline, since each stage of the pipeline is processing instructions and the slowest stage is the only thing preventing the other stages from processing even more instructions. Thus, so long as there are a lot of instructions (the process runs for a long time), and these instructions are not discarded due to other instructions, the time between each finished instruction will tend to 10ns, and the total number of instructions executed per second will tend to be $1s / 10$ ns = $1\,000\,000\,000$ ns / 10ns = 100 000 000 instructions per second, total.

Question 2: Benefits of Virtual Machines

Virtual machines are software-run emulated machines with operating systems.

a. From a company's perspective: The biggest benefit of virtual machines is the possibility of cost savings, since it's typically cheaper to buy a single massive server (system consolidation) and then run as many virtual machines and applications as your operation requires on it, as opposed to buying a separate computer for every single computer your operation may need - especially since your operation might not know how many computers they need in the first place, or there may be wasted space or processing power on each individual computer that cannot be used where in a bigger server, that space would be usable (ex. program of size 5, individual computer of size 9, you can't fit 2 programs into that computer but with a server equal to 5 such computers, you can run 9 such programs as opposed to just 5 with the individual computers), so you might run into cost issues where you buy too many, or too little and then run into severe project delays for it.

b. From a programmer's perspective: Sometimes, as a programmer, you need to do cross-platform development, where your software needs to run on operating systems other than the one you're using. Rather than getting as many different computers as operating systems you intend your code to run on, you can just install those other operating systems as virtual machines on your computer, and then test your code on those virtual machines there, saving a lot of hassle.

c. From a regular user's perspective: Sometimes, a regular user might not be sure the program they installed is safe - the program could be a virus that hijacks their system down to the boot level, for example. With a virtual machine, it's completely contained, with no possibility that malware which damages the virtual machine can actually damage the real machine. Also, sometimes you want to use an application that only exists on a different operating system. You can use a virtual machine to run that program on your own computer, without getting a new computer.

d. From a system administrator's perspective: System administrators handle the management and configuration of all the computers under their purview on a daily basis. Virtual machines make this much easier, since rather than opening up and handling countless hundreds of different machines without VMs, they can handle a single large server and make software modifications to all the VMs on the server as needed. For example, suppose you need to install a new program. Rather than opening up every single computer and manually installing the application on every single computer, the System Administrator can simply run a script on the server to install the application on every virtual machine as needed.

Question 3: Exceptions.

a. Interrupts: A notification to the CPU that something of significance has occurred and needs to be handled as soon as possible, like knocking on the door to let an owner know that someone's outside, wishing to communicate with them. Usually used for Input/Output, hardware inputs, and other applications where input is expected, but when it occurs is unknown, and things must be done according to that input. When an interrupt occurs, the CPU must save its state, load up a fresh state corresponding to what the interrupt is all about, execute the instructions relevant to the interrupt, then load back its prior state and continue as though nothing has happened.

b. Traps: Traps are software-triggered interrupts, where in the normal execution of the code, the CPU interrupts itself. It is used to detect bad instructions and methods to switch safely to kernel mode in order to execute code which requires kernel mode to work - that is, execute a pre-defined routine in kernel mode as desired. It can be used as a safe way to utilize kernel mode.

c. Hardware Interrupts vs Traps: Hardware interrupts occur due to something other than the CPU itself, such as input/output management devices, and are asynchronous to the CPU. Traps are generated by the software the CPU is currently executing, such as a bad instruction, and follow the execution of the CPU. Hardware interrupts are completely unpredictable, since they're outside the control of the code and the CPU, whilst traps can be predictable, since they're written into or implied in the code and its execution (they are an internal call or event). Both trigger the same interrupt hardware to manage interrupt problems, such as the interrupt vector table. Both save the prior state of the CPU, put the CPU in kernel mode, invoke a kernel routine, and then restore the original state of the CPU when they are done. Traps are usually utilized to execute specific code in kernel mode, whilst hardware interrupts are utilized to communicate between the hardware and a CPU waiting on the hardware. Traps are also usually intended, since unintended ones are usually called exceptions.

d. Kernel Mode vs User Mode: Interrupts are handled in kernel mode rather than user mode for many reasons. Firstly, in user mode, one cannot guarantee that the running user program is in any state to be aware that it received an interrupt and needs to handle the interrupt now - there's no reason that user mode must pay attention, unlike kernel mode, which is always connected to the hardware. Second, user-mode on most OSes does not even have permissions to touch the hardware in the first place, for security reasons, thus one requires the mode which actually has the permission - kernel mode - to handle the interrupt for the program.

Question 4: Word Count

Assignment 1: John Ming Ngo, 30020834

a. The outputs of my time command are as follows:

```
john.ngo@csx:~/457$ time ./simple_wc < a-tale-of-two-cities.txt
```

```
16272 138883 804335
```

```
real 0m1.419s
```

```
user 0m0.384s
```

```
sys 0m1.024s
```

```
john.ngo@csx:~/457$ time wc < a-tale-of-two-cities.txt
```

```
16272 138883 804335
```

```
real 0m0.018s
```

```
user 0m0.017s
```

```
sys 0m0.001s
```

b. User is the amount of time the program spent in user mode, whilst sys is the amount of time the program spent waiting on kernel mode. Based off that information, we can see that the C++ program spent 0.384 seconds in user mode, and 1.024 seconds in kernel mode. Both are vastly greater than the wc command, which used only 0.017 second and 0.001 seconds in user and system mode, respectively.

c. The biggest difference is in kernel mode - the wc command is significantly faster because it needed to spend next to no time in kernel mode at all, saving a full second there. Meanwhile, based off user mode performance, the WC command's algorithm is also likely more efficient, taking less steps to achieve the same things. The kernel mode differences can be explained (as hinted to in the lecture, and show in a later question) as the difference in the number of system calls needed to record the information, where simple_wc reads one byte of the input at a time, and so wastes a lot of time trapping, waiting on kernel mode to read the byte, then doing that over and over again, and the wc algorithm does not. Similarly, since the WC algorithm calls on the system call to read input less often, it likely requires less iterations to process everything, hence the lesser time spent in user mode.

Question 5: Rewrite the program.

See attached file/file in the same folder.

Question 6:

Here are my terminal results when testing my program versus the native wc command:

Welcome to the Department of Computer Science, University of Calgary

Assignment 1: John Ming Ngo, 30020834

This system is for use by authorized users only.

DEPARTMENT WEBPAGE: <https://ucalgary.ca/cpsc>

NEWS & ANNOUNCEMENTS: <https://ucalgary.ca/cpsc/news>

TECH SUPPORT: <https://ucalgary.ca/cpsc/tech>

Email: scihelp@ucalgary.ca

Help Desk: MS 151

Please send bugs reports and package requests to scihelp@ucalgary.ca

Last login: Tue May 19 18:10:08 2020

```
john.ngo@csx:~$ cd 457
```

```
john.ngo@csx:~/457$ ls
```

```
A1.pdf          README.txt
```

```
Assignment1_30020834.txt  Report.pdf
```

```
Assignment1.docx        romeo-and-juliet.txt
```

```
a-tale-of-two-cities.txt  simple_wc.cpp
```

```
bad_simple_wc_with_streams.cpp  smalltest.txt
```

```
makeNullByteTest.cpp      Submission_30020834.zip
```

```
myWc.cpp              'Table of Contents.html'
```

```
john.ngo@csx:~/457$ chmod +rwx * -R
```

```
john.ngo@csx:~/457$ ls
```

```
A1.pdf          README.txt
```

```
Assignment1_30020834.txt  Report.pdf
```

```
Assignment1.docx        romeo-and-juliet.txt
```

```
a-tale-of-two-cities.txt  simple_wc.cpp
```

```
bad_simple_wc_with_streams.cpp  smalltest.txt
```

```
makeNullByteTest.cpp      Submission_30020834.zip
```

```
myWc.cpp              'Table of Contents.html'
```

```
john.ngo@csx:~/457$ g++ -o wcCompiled myWc.cpp
```

```
john.ngo@csx:~/457$ g++ -o myWc myWc.cpp
```

```
john.ngo@csx:~/457$ printf "Testing a\00 null \00 example." > nullTest.txt
```

```
john.ngo@csx:~/457$ wc < nullTest.txt
```

Assignment 1: John Ming Ngo, 30020834

```
0 4 26

john.ngo@csx:~/457$ ./my-bash: _xspecs: bad array subscript
john.ngo@csx:~/457$ rm myWc
john.ngo@csx:~/457$ rm wcCompiled
john.ngo@csx:~/457$ ls
A1.pdf          Assignment1.docx  bad_simple_wc_with_streams.cpp
myWc.cpp        README.txt       romeo-and-juliet.txt  smalltest.txt      'Table of Contents.html'
Assignment1_30020834.txt  a-tale-of-two-cities.txt  makeNullByteTest.cpp
nullTest.txt  Report.pdf  simple_wc.cpp      Submission_30020834.zip

john.ngo@csx:~/457$ chmod +rwx nullTest.txt
john.ngo@csx:~/457$ rm nullText.txt
rm: cannot remove 'nullText.txt': No such file or directory
john.ngo@csx:~/457$ rm nullTest.txt
john.ngo@csx:~/457$ clear
john.ngo@csx:~/457$ ls
A1.pdf          Assignment1.docx  bad_simple_wc_with_streams.cpp
myWc.cpp        Report.pdf       simple_wc.cpp  Submission_30020834.zip
Assignment1_30020834.txt  a-tale-of-two-cities.txt  makeNullByteTest.cpp
README.txt  romeo-and-juliet.txt  smalltest.txt  'Table of Contents.html'

john.ngo@csx:~/457$ g++ -o myWc myWc.cpp
john.ngo@csx:~/457$ printf "Testing\00 U\00ser based \00\00\00 inserted null values" >
nullTest.txt

john.ngo@csx:~/457$ ls
A1.pdf          Assignment1.docx  bad_simple_wc_with_streams.cpp  myWc
nullTest.txt  Report.pdf       simple_wc.cpp  Submission_30020834.zip
Assignment1_30020834.txt  a-tale-of-two-cities.txt  makeNullByteTest.cpp
myWc.cpp  README.txt  romeo-and-juliet.txt  smalltest.txt  'Table of Contents.html'

john.ngo@csx:~/457$ wc < nullTest.txt
0 6 45

john.ngo@csx:~/457$ ./myWc < nullTest.txt
0 7 45

john.ngo@csx:~/457$ printf "It appears the difference here is if the null-byte center word
is a word at all. Doesn't really matter."

It appears the difference here is if the null-byte center word is a word at all. Doesn't really
matter.john.ngo@csx:~/457$
john.ngo@csx:~/457$ time wc < nullTest.txt
0 6 45
```

Assignment 1: John Ming Ngo, 30020834

```
real 0m0.002s
user 0m0.001s
sys 0m0.000s
john.ngo@csx:~/457$ time ./myWc < nullTest.txt
0 7 45
```

```
real 0m0.003s
user 0m0.001s
sys 0m0.001s
john.ngo@csx:~/457$ time wc < romeo-and-juliet.txt
4853 28983 178983
```

```
real 0m0.009s
user 0m0.005s
sys 0m0.001s
john.ngo@csx:~/457$ time ./m
makeNullByteTest.cpp myWc myWc.cpp
john.ngo@csx:~/457$ time ./myWc romeo-and-juliet.txt
^C
real 0m6.975s
user 0m0.001s
sys 0m0.001s
```

```
john.ngo@csx:~/457$ time ./myWc < romeo-and-juliet.txt
4853 28983 178983
```

```
real 0m0.005s
user 0m0.004s
sys 0m0.001s
john.ngo@csx:~/457$ say Small mistake there
-bash: say: command not found
john.ngo@csx:~/457$ time wc < a-tale-of-two-cities.txt
```

Assignment 1: John Ming Ngo, 30020834

16272 138883 804335

real 0m0.023s

user 0m0.018s

sys 0m0.000s

john.ngo@csx:~/457\$ time ./myWc < a-tale-of-two-cities.txt

16272 138883 804335

real 0m0.013s

user 0m0.011s

sys 0m0.002s

john.ngo@csx:~/457\$ strace -c wc < nullTest.txt

0 6 45

% time	seconds	usecs/call	calls	errors	syscall
--------	---------	------------	-------	--------	---------

43.27	0.000318	8	36	19	openat
-------	----------	---	----	----	--------

17.55	0.000129	6	19		fstat
-------	----------	---	----	--	-------

14.15	0.000104	5	18		mmap
-------	----------	---	----	--	------

11.16	0.000082	4	20		close
-------	----------	---	----	--	-------

3.95	0.000029	7	4		mprotect
------	----------	---	---	--	----------

3.13	0.000023	4	5		read
------	----------	---	---	--	------

2.18	0.000016	16	1		munmap
------	----------	----	---	--	--------

2.04	0.000015	3	4		brk
------	----------	---	---	--	-----

1.36	0.000010	10	1		write
------	----------	----	---	--	-------

0.68	0.000005	5	1		fadvise64
------	----------	---	---	--	-----------

0.54	0.000004	4	1		arch_prctl
------	----------	---	---	--	------------

0.00	0.000000	0	1	1	access
------	----------	---	---	---	--------

0.00	0.000000	0	1		execve
------	----------	---	---	--	--------

100.00	0.000735		112	20	total
--------	----------	--	-----	----	-------

john.ngo@csx:~/457\$ strace -c ./myWc < nullTest.txt

0 7 45

% time	seconds	usecs/call	calls	errors	syscall
--------	---------	------------	-------	--------	---------

35.98	0.000118	8	14	mmap	
28.35	0.000093	9	10	mprotect	
10.06	0.000033	6	5	openat	
8.84	0.000029	5	5	read	
5.18	0.000017	3	5	close	
4.88	0.000016	2	6	fstat	
3.35	0.000011	11	1	munmap	
2.13	0.000007	2	3	brk	
1.22	0.000004	4	1	arch_prctl	
0.00	0.000000	0	1	write	
0.00	0.000000	0	1	1 access	
0.00	0.000000	0	1	execve	

100.00	0.000328		53	1 total	
john.ngo@csx:~/457\$ strace -c wc < romeo-and-juliet.txt					
4853 28983 178983					
% time	seconds	usecs/call	calls	errors	syscall

42.74	0.000427	11	36	19	openat
16.52	0.000165	9	18		mmap
13.41	0.000134	6	20		close
13.41	0.000134	7	19		fstat
10.71	0.000107	7	15		read
2.10	0.000021	21	1		write
1.10	0.000011	11	1		fadvise64
0.00	0.000000	0	4		mprotect
0.00	0.000000	0	1		munmap
0.00	0.000000	0	4		brk
0.00	0.000000	0	1	1	access
0.00	0.000000	0	1		execve
0.00	0.000000	0	1		arch_prctl

Assignment 1: John Ming Ngo, 30020834

```
100.00  0.000999          122    20 total
john.ngo@csx:~/457$ strace -c ./myWc < romeo-and-juliet.txt
4853  28983 178983
% time   seconds  usecs/call   calls   errors syscall
-----
0.00  0.000000      0     5      read
0.00  0.000000      0     1      write
0.00  0.000000      0     5      close
0.00  0.000000      0     6      fstat
0.00  0.000000      0    14      mmap
0.00  0.000000      0    10      mprotect
0.00  0.000000      0     1      munmap
0.00  0.000000      0     3      brk
0.00  0.000000      0     1    1 access
0.00  0.000000      0     1      execve
0.00  0.000000      0     1      arch_prctl
0.00  0.000000      0     5      openat
```

```
100.00  0.000000          53     1 total
john.ngo@csx:~/457$ strace -c wc < a-tale-of-two-cities.txt
16272 138883 804335
% time   seconds  usecs/call   calls   errors syscall
-----
32.82  0.000576     16    36    19 openat
20.68  0.000363      6    54      read
15.21  0.000267     14    18      mmap
11.34  0.000199     10    19      fstat
9.06   0.000159      7    20      close
4.50   0.000079     19     4      mprotect
2.45   0.000043     10     4      brk
1.48   0.000026     26     1      munmap
0.80   0.000014     14     1    1 access
0.63   0.000011     11     1      execve
```

Assignment 1: John Ming Ngo, 30020834

```

0.57 0.000010    10    1    fadvise64
0.46 0.000008     8    1    arch_prctl
0.00 0.000000     0    1    write
-----
100.00 0.001755          161    20 total
john.ngo@csx:~/457$ strace -c ./myWc < a-tale-of-two-cities.txt
16272 138883 804335
% time    seconds usecs/call   calls   errors syscall
-----
32.58 0.000202     14    14     mmap
25.00 0.000155     15    10     mprotect
13.39 0.000083     16     5     openat
7.58 0.000047      9     5     read
6.61 0.000041      8     5     close
5.81 0.000036      6     6     fstat
2.90 0.000018     18     1     munmap
2.10 0.000013     13     1     1 access
1.61 0.000010     10     1     execve
1.29 0.000008      8     1     arch_prctl
1.13 0.000007      2     3     brk
0.00 0.000000      0     1     write
-----
100.00 0.000620          53     1 total
john.ngo@csx:~/457$ These are all the results.
-bash: These: command not found
john.ngo@csx:~/457$

```

a. Based off the timing results, when compared to the results for simple_wc from before, myWc is SIGNIFICANTLY faster than simple_wc. Let us investigate the strace -c of simple_wc:

```

john.ngo@csx:~/457$ strace -c ./simple_wc < romeo-and-juliet.txt
4853 28983 178983
% time    seconds usecs/call   calls   errors syscall
-----

```

Assignment 1: John Ming Ngo, 30020834

99.97	1.839927	10	178988	read
0.01	0.000197	14	14	mmap
0.01	0.000116	11	10	mprotect
0.00	0.000061	10	6	fstat
0.00	0.000055	11	5	openat
0.00	0.000044	8	5	close
0.00	0.000032	32	1	write
0.00	0.000019	19	1	1 access
0.00	0.000013	13	1	arch_prctl
0.00	0.000007	2	3	brk
0.00	0.000000	0	1	munmap
0.00	0.000000	0	1	execve

100.00	1.840471		179036	1 total
--------	----------	--	--------	---------

john.ngo@csx:~/457\$

Contrast, for myWc:

john.ngo@csx:~/457\$ strace -c ./myWc < romeo-and-juliet.txt

4853 28983 178983

% time	seconds	usecs/call	calls	errors	syscall
--------	---------	------------	-------	--------	---------

0.00	0.000000	0	5		read
0.00	0.000000	0	1		write
0.00	0.000000	0	5		close
0.00	0.000000	0	6		fstat
0.00	0.000000	0	14		mmap
0.00	0.000000	0	10		mprotect
0.00	0.000000	0	1		munmap
0.00	0.000000	0	3		brk
0.00	0.000000	0	1		1 access
0.00	0.000000	0	1		execve
0.00	0.000000	0	1		arch_prctl
0.00	0.000000	0	5		openat

100.00 0.000000 53 1 total

Consider how in a prior question, I noted that `simple_wc` takes a lot of time due to a large amount of time spent in kernel mode, and further its algorithm is less efficient, so it even spends more time in user mode. Here, comparing it to the `strace -c` of `myWc`, we can clearly see that it utilizes the `read()` call vastly more than `myWc` does, and as the `read` call is a kernel call, putting the program into kernel mode and waiting until it reads some input data to resume the program, this alone explains why `myWc` is vastly faster. `myWc` utilizes orders of magnitude less kernel calls, which means much, much less time required to wait for all the times where the CPU state is saved, waiting until other kernel calls are done and it's their turn, and then doing their kernel call, only to do it all over again later in the program.

b. Based off the 'time' results above, `myWc` is within a similar order of magnitude of time as the `wc` command, and in fact is about 50% faster than the native `wc` command. Based off the trace results above, this appears to be primarily due to significantly less kernel calls, also resulting in less user time processes needed to handle more kernel calls.

This similar order of magnitude is likely because my word count program utilizes relatively few kernel calls - from the `strace` results above, it's clear that my program uses a roughly comparable number of kernel calls, as opposed to the several orders of magnitude more kernel calls that `simple_wc` utilized, and as kernel calls are the primary source of time inefficiency from before, reducing it down to a number comparable to the native `wc` has produced similar results, especially since it appears the actual user portion of the code runs fairly fast.

To attempt to definitively state whether my program is faster or slower, we can do multiple trials:

Doing multiple trials:

Wc:

```
john.ngo@csx:~/457$ time wc < a-tale-of-two-cities.txt
```

```
16272 138883 804335
```

```
real 0m0.018s
```

```
user 0m0.016s
```

```
sys 0m0.001s
```

```
john.ngo@csx:~/457$ time wc < a-tale-of-two-cities.txt
```

```
16272 138883 804335
```

```
real 0m0.018s
```

```
user 0m0.016s
```

Assignment 1: John Ming Ngo, 30020834

```
sys 0m0.002s
john.ngo@csx:~/457$ time wc < a-tale-of-two-cities.txt
16272 138883 804335
```

```
real 0m0.018s
user 0m0.016s
sys 0m0.002s
john.ngo@csx:~/457$ time wc < a-tale-of-two-cities.txt
16272 138883 804335
```

```
real 0m0.017s
user 0m0.016s
sys 0m0.001s
john.ngo@csx:~/457$ time wc < a-tale-of-two-cities.txt
16272 138883 804335
```

```
real 0m0.018s
user 0m0.016s
sys 0m0.002s
```

myWc:

```
john.ngo@csx:~/457$ time ./myWc < a-tale-of-two-cities.txt
16272 138883 804335
```

```
real 0m0.014s
user 0m0.010s
sys 0m0.003s
john.ngo@csx:~/457$ time ./myWc < a-tale-of-two-cities.txt
16272 138883 804335
```

```
real 0m0.013s
user 0m0.012s
sys 0m0.001s
john.ngo@csx:~/457$ time ./myWc < a-tale-of-two-cities.txt
```

Assignment 1: John Ming Ngo, 30020834

```
16272 138883 804335
```

```
real 0m0.013s
```

```
user 0m0.011s
```

```
sys 0m0.002s
```

```
john.ngo@csx:~/457$ time ./myWc < a-tale-of-two-cities.txt
```

```
16272 138883 804335
```

```
real 0m0.013s
```

```
user 0m0.012s
```

```
sys 0m0.000s
```

```
john.ngo@csx:~/457$ time ./myWc < a-tale-of-two-cities.txt
```

```
16272 138883 804335
```

```
real 0m0.014s
```

```
user 0m0.012s
```

```
sys 0m0.002s
```

```
john.ngo@csx:~/457$
```

Eyeballing these numbers, based off the time average of these values, myWc is definitively faster for a relatively large input file, such as 'a tale of two cities', due primarily due to significantly less user processing time despite surprisingly similar amounts of kernel time.

Screenshots of the above here:

Assignment 1: John Ming Ngo, 30020834

```
johnngo@csc:~/.457
Using username "johnngo".
johnngo@linux.cpsc.ucalgary.ca's password:
Welcome to the Department of Computer Science, University of Calgary
This system is for use by authorized users only.

DEPARTMENT WEBPAGE: https://ucalgary.ca/cpsc
NEWS & ANNOUNCEMENTS: https://ucalgary.ca/cpsc/news
TECH SUPPORT: https://ucalgary.ca/cpsc/tech
Email: scihelp@ucalgary.ca
Help Desk: M3 151

Please send bugs reports and package requests to scihelp@ucalgary.ca

Last login: Tue May 19 18:10:08 2020
johnngo@csc:~$ cd 457
johnngo@csc:~/457$ ls
Al.pdf          README.txt
Assignment1_30020834.txt  Report.pdf
Assignment1.docx  romeo-and-juliet.txt
a-tale-of-two-cities.txt  simple_wc.cpp
bad simple wc with streams.cpp  smallTest.txt
makeNullByteTest.cpp  Submission_30020834.zip
myWc.cpp          'Table of Contents.html'
johnngo@csc:~/457$ chmod +rxw * -R
johnngo@csc:~/457$ ls
Al.pdf          README.txt
Assignment1_30020834.txt  Report.pdf
Assignment1.docx  romeo-and-juliet.txt
a-tale-of-two-cities.txt  simple_wc.cpp
bad simple wc with streams.cpp  smallTest.txt
makeNullByteTest.cpp  Submission_30020834.zip
myWc.cpp          'Table of Contents.html'
johnngo@csc:~/457$ g++ -o wcCompiled myWc.cpp
johnngo@csc:~/457$ g++ -o myWc myWc.cpp
johnngo@csc:~/457$ printf "Testing a\000 null \000 example." > nullTest.txt
johnngo@csc:~/457$ wc < nullTest.txt
0  4 26
johnngo@csc:~/457$ ./my-bash: _xspecs: bad array subscript
johnngo@csc:~/457$ rm myWc
johnngo@csc:~/457$ rm wcCompiled
johnngo@csc:~/457$ ls
Al.pdf          Assignment1.docx      bad simple wc with streams.cpp  myWc.cpp  README.txt  romeo-and-juliet.txt  smallTest.txt  'Table of Contents.html'
Assignment1_30020834.txt  a-tale-of-two-cities.txt  makeNullByteTest.cpp          nullTest.txt  Report.pdf  simple_wc.cpp  Submission_30020834.zip
johnngo@csc:~/457$ chmod +rxw nullTest.txt
johnngo@csc:~/457$ rm nullTest.txt
rm: cannot remove 'nullTest.txt': No such file or directory
johnngo@csc:~/457$ rm nullTest.txt
johnngo@csc:~/457$ clear
```

```
johnngo@csc:~/457$ clear
johnngo@csc:~/457$ ls
Al.pdf          Assignment1.docx      bad simple wc with streams.cpp  myWc.cpp  README.txt  romeo-and-juliet.txt  smallTest.txt  'Table of Contents.html'
Assignment1_30020834.txt  a-tale-of-two-cities.txt  makeNullByteTest.cpp          nullTest.txt  Report.pdf  simple_wc.cpp  Submission_30020834.zip
johnngo@csc:~/457$ g++ -o myWc myWc.cpp
johnngo@csc:~/457$ printf "Testing\000 U\000ser based \000\000\000 inserted null values" > nullTest.txt
johnngo@csc:~/457$ ls
Al.pdf          Assignment1.docx      bad simple wc with streams.cpp  myWc.cpp  nullTest.txt  README.txt  romeo-and-juliet.txt  smallTest.txt  'Table of Contents.html'
Assignment1_30020834.txt  a-tale-of-two-cities.txt  makeNullByteTest.cpp          myWc.cpp  nullTest.txt  Report.pdf  simple_wc.cpp  Submission_30020834.zip
johnngo@csc:~/457$ wc < nullTest.txt
0  6 45
johnngo@csc:~/457$ ./myWc < nullTest.txt
0  7  45
johnngo@csc:~/457$ printf "It appears the difference here is if the null-byte center word is a word at all. Doesn't really matter."
It appears the difference here is if the null-byte center word is a word at all. Doesn't really matter.johnngo@csc:~/457$
johnngo@csc:~/457$ time wc < nullTest.txt
0  6 45

real    0m0.002s
user    0m0.001s
sys     0m0.000s
johnngo@csc:~/457$ time ./myWc < nullTest.txt
0  7  45

real    0m0.003s
user    0m0.001s
sys     0m0.001s
johnngo@csc:~/457$ time wc < romeo-and-juliet.txt
4853 28983 178983

real    0m0.009s
user    0m0.005s
sys     0m0.001s
johnngo@csc:~/457$ time ./m
makeNullByteTest.cpp myWc          myWc.cpp
johnngo@csc:~/457$ time ./myWc romeo-and-juliet.txt
^C
real    0m6.975s
user    0m0.001s
sys     0m0.001s
johnngo@csc:~/457$ time ./myWc < romeo-and-juliet.txt
4853 28983 178983

real    0m0.005s
user    0m0.004s
sys     0m0.001s
johnngo@csc:~/457$ say Small mistake there
-bash: say: command not found
```

Assignment 1: John Ming Ngo, 30020834

```
johnngo@csc:~/.457
~bash: say: command not found
johnngo@csc:~/.457$ time wc < a-tale-of-two-cities.txt
16272 138883 804335

real    0m0.023s
user    0m0.018s
sys     0m0.000s
johnngo@csc:~/.457$ time ./myWc < a-tale-of-two-cities.txt
16272 138883 804335

real    0m0.013s
user    0m0.011s
sys     0m0.002s
johnngo@csc:~/.457$ strace -c wc < nullTest.txt
0 0 45
% time    seconds  usecs/call   calls   errors syscall
-----
43.27 0.000318      8      36      19 openat
17.55 0.000129      6      19       0 fstat
14.15 0.000104      5      18       0 mmap
11.16 0.000082      4      20       0 close
3.95 0.000029      7       4       0 mprotect
3.13 0.000023      4       5       0 read
2.18 0.000016     16       1       0 munmap
2.04 0.000015      3       4       0 brk
1.36 0.000010     10       1       0 write
0.68 0.000005      5       1       0 fadvise64
0.54 0.000004      4       1       0 arch_prctl
0.00 0.000000      0       1       0 l access
0.00 0.000000      0       1       0 execve
-----
100.00 0.000735     112     20 total
johnngo@csc:~/.457$ strace -c ./myWc < nullTest.txt
0 7 45
% time    seconds  usecs/call   calls   errors syscall
-----
35.98 0.000118      8      14       0 mmap
28.35 0.000093      9      10       0 mprotect
10.06 0.000033      6       5       0 openat
8.84 0.000029      5       5       0 read
5.18 0.000017      3       5       0 close
4.88 0.000016      2       6       0 fstat
3.35 0.000011     11       1       0 munmap
2.13 0.000007      2       3       0 brk
1.22 0.000004      4       1       0 arch_prctl
0.00 0.000000      0       1       0 write
0.00 0.000000      0       1       0 l access
0.00 0.000000      0       1       0 execve
-----
100.00 0.000999     122     20 total
johnngo@csc:~/.457$ strace -c ./myWc < romeo-and-juliet.txt
4853 28983 178983
% time    seconds  usecs/call   calls   errors syscall
-----
42.74 0.000427     11      36      19 openat
16.52 0.000165      9      18       0 mmap
13.41 0.000134      6      20       0 close
13.41 0.000134      7      19       0 fstat
10.71 0.000107      7      15       0 read
2.10 0.000021     21       1       0 write
1.10 0.000011     11       1       0 fadvise64
0.00 0.000000      0       4       0 mprotect
0.00 0.000000      0       1       0 munmap
0.00 0.000000      0       4       0 brk
0.00 0.000000      0       1       0 l access
0.00 0.000000      0       1       0 execve
0.00 0.000000      0       1       0 arch_prctl
-----
100.00 0.000999     122     20 total
johnngo@csc:~/.457$ strace -c ./myWc < romeo-and-juliet.txt
4853 28983 178983
% time    seconds  usecs/call   calls   errors syscall
-----
0.00 0.000000      0       5       0 read
0.00 0.000000      0       1       0 write
0.00 0.000000      0       5       0 close
0.00 0.000000      0       6       0 fstat
0.00 0.000000      0      14       0 mmap
0.00 0.000000      0      10       0 mprotect
0.00 0.000000      0       1       0 munmap
0.00 0.000000      0       3       0 brk
0.00 0.000000      0       1       0 l access
0.00 0.000000      0       1       0 execve
0.00 0.000000      0       1       0 arch_prctl
0.00 0.000000      0       5       0 openat
-----
100.00 0.000000      53       1 total
johnngo@csc:~/.457$ strace -c wc < a-tale-of-two-cities.txt
16272 138883 804335
% time    seconds  usecs/call   calls   errors syscall
-----
32.82 0.000576     16      36      19 openat
20.68 0.000363      6      54       0 read
15.21 0.000267     14      18       0 mmap
11.24 0.000159     10      19       0 fstat
9.06 0.000159      7      20       0 close
4.50 0.000079     19       4       0 mprotect
```


Assignment 1: John Ming Ngo, 30020834

```
john.ngo@csx:~/457
-----
32.82 0.000576 16 36 19 openat
20.68 0.000363 6 54 read
15.21 0.000267 14 18 mmap
11.34 0.000199 10 19 fstat
9.06 0.000159 7 20 close
4.50 0.000079 19 4 mprotect
2.45 0.000043 10 4 brk
1.48 0.000026 26 1 munmap
0.80 0.000014 14 1 l access
0.63 0.000011 11 1 execve
0.57 0.000010 10 1 fadvise64
0.46 0.000008 8 1 arch_prctl
0.00 0.000000 0 1 write
-----
100.00 0.001755 161 20 total
john.ngo@csx:~/457$ strace -c ./myWc < a-tale-of-two-cities.txt
16272 138883 804335
% time seconds usecs/call calls errors syscall
-----
32.58 0.000202 14 14 mmap
25.00 0.000155 15 10 mprotect
13.39 0.000083 16 5 openat
7.58 0.000047 9 5 read
6.61 0.000041 8 5 close
5.81 0.000036 6 6 fstat
2.90 0.000018 18 1 munmap
2.10 0.000013 13 1 l access
1.61 0.000010 10 1 execve
1.29 0.000008 8 1 arch_prctl
1.13 0.000007 2 3 brk
0.00 0.000000 0 1 write
-----
100.00 0.000620 53 1 total
john.ngo@csx:~/457$ These are all the results.
-bash: These: command not found
john.ngo@csx:~/457$ ^C
john.ngo@csx:~/457$ time wc < a-tale-of-two-cities.txt
16272 138883 804335

real 0m0.019s
user 0m0.015s
sys 0m0.002s
john.ngo@csx:~/457$ time wc < a-tale-of-two-cities.txt
16272 138883 804335

real 0m0.017s
user 0m0.015s
sys 0m0.002s

-bash: These: command not found
john.ngo@csx:~/457$ ^C
john.ngo@csx:~/457$ time wc < a-tale-of-two-cities.txt
16272 138883 804335

real 0m0.019s
user 0m0.015s
sys 0m0.002s
john.ngo@csx:~/457$ time wc < a-tale-of-two-cities.txt
16272 138883 804335

real 0m0.017s
user 0m0.015s
sys 0m0.002s
john.ngo@csx:~/457$ time wc < a-tale-of-two-cities.txt
16272 138883 804335

real 0m0.018s
user 0m0.016s
sys 0m0.002s
john.ngo@csx:~/457$ time wc < a-tale-of-two-cities.txt
16272 138883 804335

real 0m0.019s
user 0m0.019s
sys 0m0.000s
john.ngo@csx:~/457$ time wc < a-tale-of-two-cities.txt
16272 138883 804335

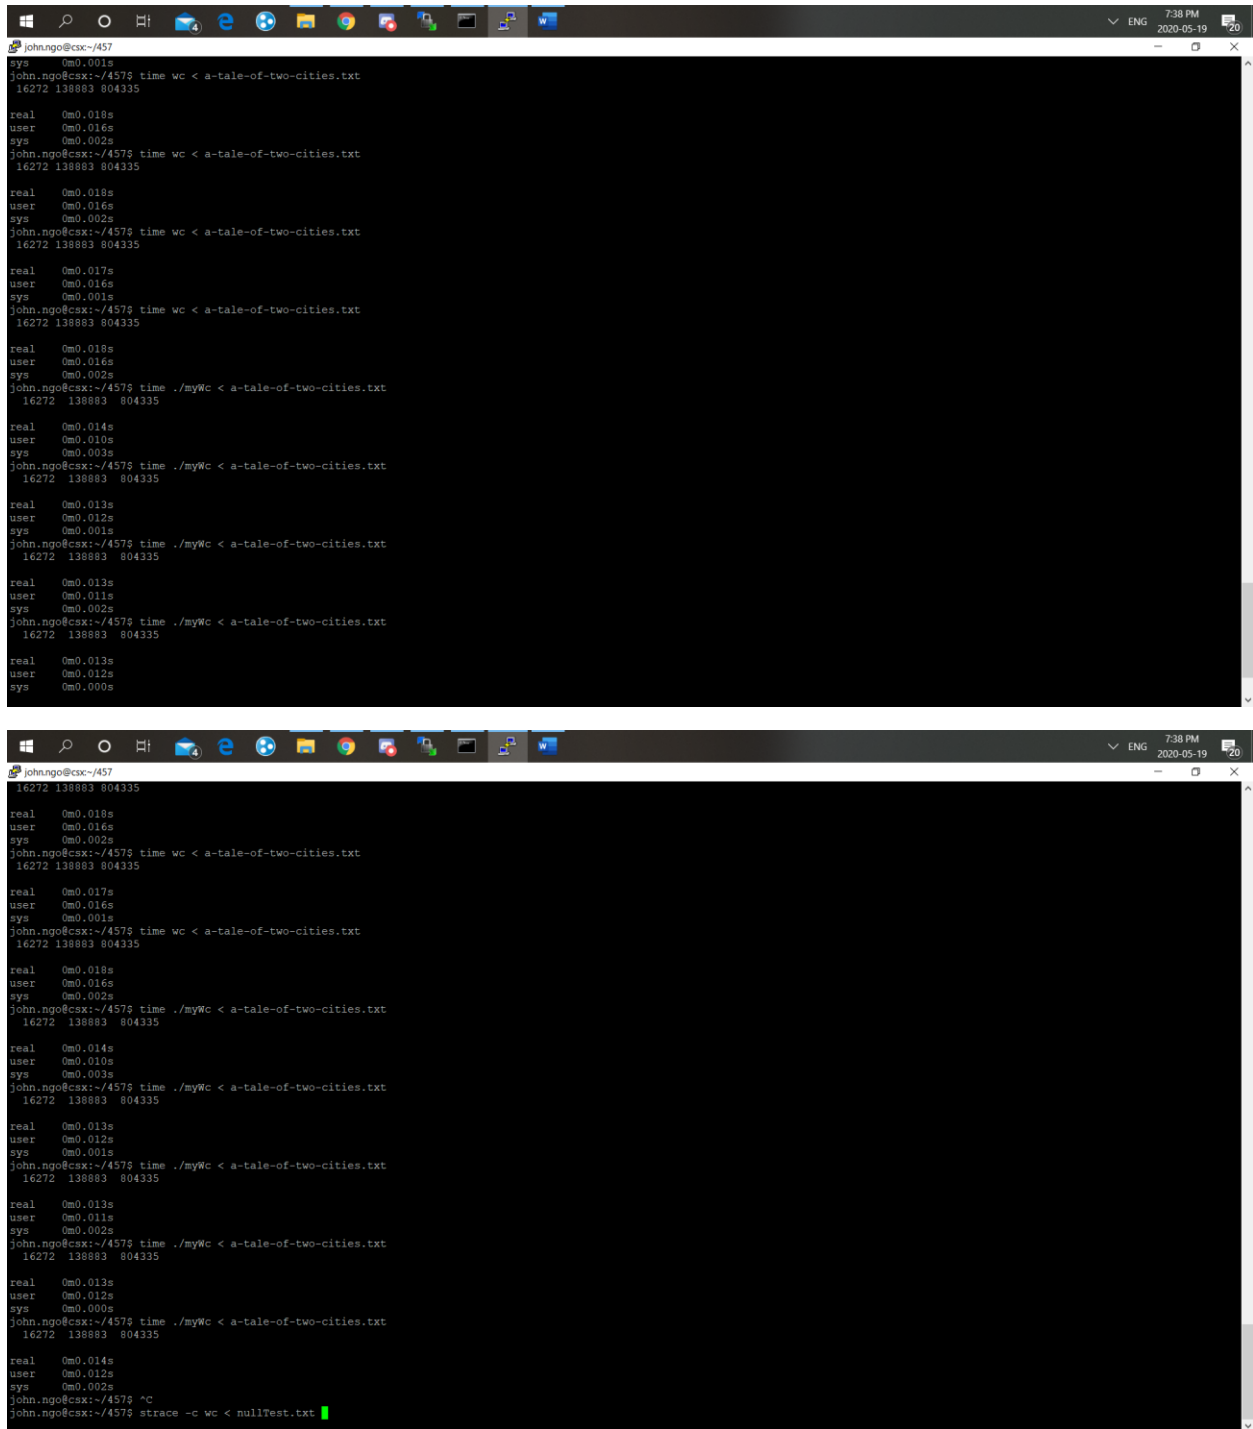
real 0m0.019s
user 0m0.017s
sys 0m0.001s
john.ngo@csx:~/457$ time ./Wc < a-tale-of-two-cities.txt
-bash: ./Wc: No such file or directory

real 0m0.002s
user 0m0.001s
sys 0m0.000s
john.ngo@csx:~/457$ time wc < a-tale-of-two-cities.txt
16272 138883 804335

real 0m0.018s
user 0m0.016s
sys 0m0.001s
john.ngo@csx:~/457$ time wc < a-tale-of-two-cities.txt
16272 138883 804335

real 0m0.018s
user 0m0.016s
```

Assignment 1: John Ming Ngo, 30020834



```
johnngo@csc:~/457
sys 0m0.001s
johnngo@csc:~/457$ time wc < a-tale-of-two-cities.txt
16272 138883 804335

real 0m0.018s
user 0m0.016s
sys 0m0.002s
johnngo@csc:~/457$ time wc < a-tale-of-two-cities.txt
16272 138883 804335

real 0m0.018s
user 0m0.016s
sys 0m0.002s
johnngo@csc:~/457$ time wc < a-tale-of-two-cities.txt
16272 138883 804335

real 0m0.017s
user 0m0.016s
sys 0m0.001s
johnngo@csc:~/457$ time wc < a-tale-of-two-cities.txt
16272 138883 804335

real 0m0.018s
user 0m0.016s
sys 0m0.002s
johnngo@csc:~/457$ time ./myWc < a-tale-of-two-cities.txt
16272 138883 804335

real 0m0.014s
user 0m0.010s
sys 0m0.003s
johnngo@csc:~/457$ time ./myWc < a-tale-of-two-cities.txt
16272 138883 804335

real 0m0.013s
user 0m0.012s
sys 0m0.001s
johnngo@csc:~/457$ time ./myWc < a-tale-of-two-cities.txt
16272 138883 804335

real 0m0.013s
user 0m0.011s
sys 0m0.002s
johnngo@csc:~/457$ time ./myWc < a-tale-of-two-cities.txt
16272 138883 804335

real 0m0.013s
user 0m0.012s
sys 0m0.000s
johnngo@csc:~/457$

16272 138883 804335

real 0m0.018s
user 0m0.016s
sys 0m0.002s
johnngo@csc:~/457$ time wc < a-tale-of-two-cities.txt
16272 138883 804335

real 0m0.017s
user 0m0.016s
sys 0m0.001s
johnngo@csc:~/457$ time wc < a-tale-of-two-cities.txt
16272 138883 804335

real 0m0.018s
user 0m0.016s
sys 0m0.002s
johnngo@csc:~/457$ time ./myWc < a-tale-of-two-cities.txt
16272 138883 804335

real 0m0.014s
user 0m0.010s
sys 0m0.003s
johnngo@csc:~/457$ time ./myWc < a-tale-of-two-cities.txt
16272 138883 804335

real 0m0.013s
user 0m0.012s
sys 0m0.001s
johnngo@csc:~/457$ time ./myWc < a-tale-of-two-cities.txt
16272 138883 804335

real 0m0.013s
user 0m0.011s
sys 0m0.002s
johnngo@csc:~/457$ time ./myWc < a-tale-of-two-cities.txt
16272 138883 804335

real 0m0.013s
user 0m0.012s
sys 0m0.000s
johnngo@csc:~/457$ time ./myWc < a-tale-of-two-cities.txt
16272 138883 804335

real 0m0.014s
user 0m0.012s
sys 0m0.002s
johnngo@csc:~/457$ ^C
johnngo@csc:~/457$ strace -c wc < nullTest.txt
```